

```

1 ### Assignment 4 Generative Adversarial Nets (Unconditional, 10 pts)
2
3
4 In this exercise, we will implement a Generative Adversarial Net (GAN), specifically, a Wasserstein GAN and train it on the M
5
6 **Submit**
7 1. (<font color='red'>Doc A</font>) Include the two figures at the end in the pdf generated by the latex file with Exercise 2
8 2. (<font color='red'>Doc B</font>) The completed *.ipynb file with all the command outputs (can be created by saving the fil

```

Setup

1. In Colab, open tab Runtime > Change runtime type, choose *python3* and *T4 GPU*.
2. Run the following command to set up the environment. (Takes ~ 1.5 min)

```
1 ! pip install --quiet "ipython[notebook]==7.34.0, <8.17.0" "setuptools>=68.0.0, <68.3.0" "torch==1.13.0" "matplotlib" "torc
```

```

807.9/807.9 kB 5.2 MB/s eta 0:00:00
890.1/890.1 MB 1.2 MB/s eta 0:00:00
1.6/1.6 MB 49.0 MB/s eta 0:00:00
849.3/849.3 kB 45.7 MB/s eta 0:00:00
557.1/557.1 MB 1.9 MB/s eta 0:00:00
317.1/317.1 MB 4.5 MB/s eta 0:00:00
21.0/21.0 MB 47.8 MB/s eta 0:00:00
7.0/7.0 MB 71.8 MB/s eta 0:00:00
7.0/7.0 MB 78.3 MB/s eta 0:00:00
6.9/6.9 MB 72.8 MB/s eta 0:00:00
6.9/6.9 MB 68.8 MB/s eta 0:00:00
6.9/6.9 MB 72.5 MB/s eta 0:00:00
6.8/6.8 MB 67.2 MB/s eta 0:00:00
6.8/6.8 MB 73.3 MB/s eta 0:00:00
6.9/6.9 MB 15.1 MB/s eta 0:00:00
6.0/6.0 MB 89.2 MB/s eta 0:00:00
6.0/6.0 MB 83.6 MB/s eta 0:00:00
24.2/24.2 MB 58.5 MB/s eta 0:00:00
24.3/24.3 MB 58.5 MB/s eta 0:00:00
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is
torchaudio 2.3.1+cu121 requires torch==2.3.1, but you have torch 1.13.0 which is incompatible.
torchtext 0.18.0 requires torch>=2.3.0, but you have torch 1.13.0 which is incompatible.

```

Let's start with importing our standard set of libraries.

```

1 import torch
2 from torch import nn, optim, autograd
3 import torchvision
4 import torchvision.transforms as transforms
5 import matplotlib.pyplot as plt
6 import torchvision.utils as vutils
7 from dataclasses import dataclass
8 import time
9 import sys
10 %matplotlib inline
11 torch.set_num_threads(1)
12 torch.manual_seed(1)
13
14
15 device = torch.device("cuda:0") if torch.cuda.is_available() else torch.device("cpu")
16
17 if device == torch.device("cuda:0"):
18     print('Everything looks good; continue')
19 else:
20     # It is OK if you cannot connect to a GPU. In this case, training the model for
21     # 2 epoch is sufficient to get full mark. (NOTE THAT 2 epoch takes approximately 1.5 hours to train for CPU)
22     print('GPU is not detected. Make sure you have chosen the right runtime type')

```

```
Everything looks good; continue
```

Dataloaders and hyperparameters (0 pt)

```

1 @dataclass
2 class Hyperparameter:
3     batchsize: int         = 64
4     num_epochs: int        = 5
5     latent_size: int       = 32
6     n_critic: int          = 5
7     critic_size: int       = 1024
8     generator_size: int    = 1024
9     critic_hidden_size: int = 1024
10    gp_lambda: float       = 10.
11
12 hp = Hyperparameter()
13
14 transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5,), (0.5,))])
15
16 dataset = torchvision.datasets.MNIST("mnist", download=True, transform=transform)
17 dataloader = torch.utils.data.DataLoader(dataset, batch_size=hp.batchsize, num_workers=1, shuffle=True, drop_last=True, pin_n

```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz>
 Failed to download (trying next):
 HTTP Error 403: Forbidden

Downloading <https://oss-ci-datasets.s3.amazonaws.com/mnist/train-images-idx3-ubyte>
 Downloading <https://oss-ci-datasets.s3.amazonaws.com/mnist/train-images-idx3-ubyte>
 100% 9912422/9912422 [00:00<00:00, 14971358.60it/s]
 Extracting mnist/MNIST/raw/train-images-idx3-ubyte.gz to mnist/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz>
 Failed to download (trying next):
 HTTP Error 403: Forbidden

Downloading <https://oss-ci-datasets.s3.amazonaws.com/mnist/train-labels-idx1-ubyte>
 Downloading <https://oss-ci-datasets.s3.amazonaws.com/mnist/train-labels-idx1-ubyte>
 100% 28881/28881 [00:00<00:00, 422290.41it/s]
 Extracting mnist/MNIST/raw/train-labels-idx1-ubyte.gz to mnist/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz>
 Failed to download (trying next):
 HTTP Error 403: Forbidden

Downloading <https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-images-idx3-ubyte>
 Downloading <https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-images-idx3-ubyte>
 100% 1648877/1648877 [00:00<00:00, 3850616.34it/s]
 Extracting mnist/MNIST/raw/t10k-images-idx3-ubyte.gz to mnist/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz>
 Failed to download (trying next):
 HTTP Error 403: Forbidden

Downloading <https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte>
 Downloading <https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte>
 100% 4542/4542 [00:00<00:00, 143919.87it/s]
 Extracting mnist/MNIST/raw/t10k-labels-idx1-ubyte.gz to mnist/MNIST/raw

✓ Building Models (2 pts)

After examining the preprocessing steps, we can now start building the models, including the generator for generating new images from random noise, and a critic of the realness of the image.

In this assignment we adopt the implementation of [DCGAN](#), which is a direct extension of [GAN](#), with convolutional and convolutional-transpose layers in the critic and generator, respectively. Specifically, we will use the [ConvTranspose2d](#) layers to upscale the noise.

Moreover, we apply an improved version of [Wasserstein-GAN](#) with a [Gradient Penalty](#) (you may read Algorithm 1 to fully understand the code we are implementing).

```

1 # Define the generator
2
3 class Generator(nn.Module):
4     def __init__(self):
5         super(Generator, self).__init__()
6
7         #VVVVVVVVVV TO BE COMPLETE (START) VVVVVVVVVVV
8         # Add latent embedding layer to adjust the dimension of the input (1 pt)
9
10        # Hint: you should use the hyperparameters defined above
11
12        self.latent_embedding = nn.Sequential(nn.Linear(hp.latent_size , hp.generator_size ))
13
14        # ^^^^^^^^^^^ TO BE COMPLETE (END) ^^^^^^^^^^^
15
16
17        # Transposed CNN layers to transfer noise to image
18
19        self.tcn = nn.Sequential(
20            # input is Z, going into a convolution
21            nn.ConvTranspose2d(hp.generator_size, hp.generator_size, kernel_size=4, stride=1, padding= 0),
22            nn.BatchNorm2d(hp.generator_size),
23            nn.ReLU(inplace=True),
24            # upscaling
25            nn.ConvTranspose2d(hp.generator_size, hp.generator_size // 2, 3, 2, 1),
26            nn.BatchNorm2d(hp.generator_size // 2),
27            nn.ReLU(inplace=True),
28            # upscaling
29            nn.ConvTranspose2d(hp.generator_size // 2, hp.generator_size // 4, 4, 2, 1),
30            nn.BatchNorm2d(hp.generator_size // 4),
31            nn.ReLU(inplace=True),
32            nn.ConvTranspose2d(hp.generator_size // 4, 1, 4, 2, 1),
33            nn.Tanh()
34        )
35
36
37    def forward(self, latent):
38        vec_latent = self.latent_embedding(latent).reshape(-1, hp.generator_size, 1, 1)
39        return self.tcn(vec_latent)
40
41
42 # Define the critic
43
44 class Critic(nn.Module):
45     def __init__(self):
46         super(Critic, self).__init__()
47
48        # CNN layers that perform downscaling
49        self.cnn_net = nn.Sequential(
50            nn.Conv2d(1, hp.critic_size // 4, 3, 2),
51            nn.InstanceNorm2d(hp.critic_size // 4, affine=True),
52            nn.LeakyReLU(0.2, inplace=True),
53            nn.Conv2d(hp.critic_size // 4, hp.critic_size // 2, 3, 2),
54            nn.InstanceNorm2d(hp.critic_size // 2, affine=True),
55            nn.LeakyReLU(0.2, inplace=True),
56            nn.Conv2d(hp.critic_size // 2, hp.critic_size, 3, 2),
57            nn.InstanceNorm2d(hp.critic_size, affine=True),
58            nn.LeakyReLU(0.2, inplace=True),
59            nn.Flatten(),
60        )
61
62        # Linear layers that produce the output from the features
63        self.critic_net = nn.Sequential(
64            nn.Linear(hp.critic_size * 4, hp.critic_hidden_size),
65            nn.LeakyReLU(0.2, inplace=True),
66
67            #VVVVVVVVVV TO BE COMPLETE (START) VVVVVVVVVVV
68            # Add the last layer to reflect the output (1 pt)
69
70            nn.Linear(hp.critic_hidden_size ,1)
71
72            # Hint: Given an image, the output of the critic is a value (or a scalar)
73
74            # ^^^^^^^^^^^ TO BE COMPLETE (END) ^^^^^^^^^^^
75        )
76
77    def forward(self, image):

```

```

78     cnn_features = self.cnn_net(image)
79     return self.critic_net(cnn_features)
80

```

✓ Before Training

Next we define the two models and the optimizers. We use the [AdamW](#) algorithm.

```

1 critic, generator = Critic().to(device), Generator().to(device)
2
3 critic_optimizer = optim.AdamW(critic.parameters(), lr=1e-4, betas=(0., 0.9))
4 generator_optimizer = optim.AdamW(generator.parameters(), lr=1e-4, betas=(0., 0.9))

```

✓ Training pipeline (6 points)

Finally, we perform training on the two networks. The training consists of two steps: (1) Updating discriminators for n_{critic} steps (such that we have an optimal critic): here we use an aggregation of three loss functions, (a) The real loss (the output scalar of the critic for real images); (b) The fake loss (same value for fake images); (c) The [gradient penalty](#). (2) Updating generators by only considering the fake loss (to fool the critic).

```

1 img_list, generator_losses, critic_losses = [], [], []
2 iters = 0
3 fixed_noise = torch.randn((64, hp.latent_size), device=device)
4 grad_tensor = torch.ones((hp.batchsize, 1), device=device)
5 start_time = time.time()
6
7 # ref: https://www.youtube.com/watch?v=ILpC3b-819Q
8 def loss_fn(y_pred, y_true):
9     return torch.mean(y_pred * y_true)
10
11 for epoch in range(hp.num_epochs):
12     for batch_idx, data in enumerate(data_loader, 0):
13         real_images = data[0].to(device)
14
15         real_labels = torch.ones(hp.batchsize, device=device) # real label = 1
16
17         # Update Critic
18         critic_optimizer.zero_grad()
19
20         # (a) Real loss
21         critic_output_real = critic(real_images)
22         critic_loss_real = loss_fn(critic_output_real, real_labels) # changed critic_output_real.mean() #
23
24         # (b) Fake loss
25
26         #VVVVVVVVVV TO BE COMPLETE (START) VVVVVVVVVVV
27         # Implement the fake loss
28
29         # (1) Generating a noise tensor (of dimension (batch_size, latent_size)), you are required to
30         # use the hyperparameters in the hp class (0.5 pt)
31
32         noise = torch.randn((hp.batchsize, hp.latent_size), device=device)
33
34         # (2) Generate fake images using the generator (hint: you are not supposed to perform gradient
35         # update on the generator) (1.5 pts)
36
37
38         fake_image = generator.forward(noise)
39         fake_labels = -real_labels # fake label = -1
40         flipped_fake_labels = real_labels # here, fake label = 1
41
42         # (3) Calculate the fake loss using the output of the generator (1 pt)
43         critic_output_fake = critic(fake_image.detach())
44         critic_loss_fake = loss_fn(critic_output_fake, fake_labels) # critic_output_fake.mean() #
45
46         #^^^^^^^^^^^^ TO BE COMPLETE (END) ^^^^^^^^^^^^^
47
48         # (c) Gradient penalty
49         alpha = torch.rand((hp.batchsize, 1, 1, 1), device=device)
50         interpolates = (alpha * real_images + ((1 - alpha) * fake_image)).requires_grad_(True)

```

```

51 ..... d_interpolates = critic(interpolates)
52 ..... gradients = autograd.grad(outputs=d_interpolates, inputs=interpolates, grad_outputs=grad_tensor, create_graph=True,
53 ..... gradient_penalty = hp.gp_lambda * ((gradients.view(hp.batchsize, -1).norm(dim=1) - 1.) ** 2).mean()
54
55 ..... #VVVVVVVVVV TO BE COMPLETE (START) VVVVVVVVVVV
56 ..... # Implement the aggregated loss using the above three components, be careful with the signs (1 pt)
57
58 ..... critic_loss = 0.5*(critic_loss_real + critic_loss_fake) + gradient_penalty
59
60 ..... #^^^^^^^^^^ TO BE COMPLETE (END) ^^^^^^^^^^
61
62 ..... critic_loss.backward()
63 ..... critic_optimizer.step()
64
65 ..... if batch_idx % hp.n_critic == 0:
66 .....     # Update Generator
67 .....     generator_optimizer.zero_grad()
68
69
70 ..... #VVVVVVVVVV TO BE COMPLETE (START) VVVVVVVVVVV
71 ..... # Implement the generator loss (2 pts)
72
73 ..... noise = torch.randn((hp.batchsize, hp.latent_size), device = device)
74 ..... fake_image = generator.forward(noise)
75 ..... critic_output_fake = critic(fake_image)
76 ..... generator_loss = loss_fn(critic_output_fake, flipped_fake_labels) #critic_output_fake.mean() #
77
78 ..... #^^^^^^^^^^ TO BE COMPLETE (END) ^^^^^^^^^^
79
80 ..... generator_loss.backward()
81 ..... generator_optimizer.step()
82
83 ..... # Output training stats
84 ..... if batch_idx % 100 == 0:
85 .....     elapsed_time = time.time() - start_time
86 .....     print(f"[epoch:>2]/{hp.num_epochs}][{iters:>7}][{elapsed_time:8.2f}s]\t"
87 .....           f"d_loss/g_loss: {critic_loss.item():4.2}/{generator_loss.item():4.2}\t")
88
89 ..... # Save Losses for plotting later
90 ..... generator_losses.append(generator_loss.item())
91 ..... critic_losses.append(critic_loss.item())
92
93 ..... # Check how the generator is doing by saving G's output on fixed_noise
94 ..... if (iters % 500 == 0) or ((epoch == hp.num_epochs - 1) and (batch_idx == len(dataloader) - 1)):
95 .....     with torch.no_grad(): fake_images = generator(fixed_noise).cpu()
96 .....     img_list.append(vutils.make_grid(fake_images, padding=2, normalize=True))
97
98 ..... iters += 1

```

```

[ 0/5][    0][   6.25s]    d_loss/g_loss:  2.0/0.08
[ 0/5][   100][  24.37s]    d_loss/g_loss: -1.3/-0.16
[ 0/5][   200][  42.92s]    d_loss/g_loss: -1.2/0.44
[ 0/5][   300][  62.12s]    d_loss/g_loss: -1.4/0.71
[ 0/5][   400][  81.73s]    d_loss/g_loss: -1.9/-0.35
[ 0/5][   500][ 101.16s]    d_loss/g_loss: -2.1/ 1.0
[ 0/5][   600][ 120.31s]    d_loss/g_loss: -2.1/0.47
[ 0/5][   700][ 139.48s]    d_loss/g_loss: -2.5/ 1.5
[ 0/5][   800][ 158.67s]    d_loss/g_loss: -2.1/ 1.5
[ 0/5][   900][ 178.03s]    d_loss/g_loss: -2.7/ 0.9
[ 1/5][   937][ 185.35s]    d_loss/g_loss: -2.6/ 0.7
[ 1/5][  1037][ 204.60s]    d_loss/g_loss: -2.9/ 1.1
[ 1/5][  1137][ 223.80s]    d_loss/g_loss: -3.0/ 1.4
[ 1/5][  1237][ 242.99s]    d_loss/g_loss: -3.0/ 2.2
[ 1/5][  1337][ 262.22s]    d_loss/g_loss: -3.1/ 1.0
[ 1/5][  1437][ 281.38s]    d_loss/g_loss: -2.9/ 2.3
[ 1/5][  1537][ 300.73s]    d_loss/g_loss: -3.1/ 1.9
[ 1/5][  1637][ 319.96s]    d_loss/g_loss: -3.0/ 1.9
[ 1/5][  1737][ 339.24s]    d_loss/g_loss: -0.38/-1.1
[ 1/5][  1837][ 358.49s]    d_loss/g_loss: -0.19/-0.33
[ 2/5][  1874][ 365.78s]    d_loss/g_loss: -0.43/0.21
[ 2/5][  1974][ 385.03s]    d_loss/g_loss: -0.56/0.21
[ 2/5][  2074][ 404.31s]    d_loss/g_loss: -0.35/-0.18
[ 2/5][  2174][ 423.50s]    d_loss/g_loss: -0.27/-0.44
[ 2/5][  2274][ 442.71s]    d_loss/g_loss: -0.48/-0.14
[ 2/5][  2374][ 461.92s]    d_loss/g_loss: -0.0017/0.085
[ 2/5][  2474][ 481.16s]    d_loss/g_loss: -0.17/-1.1
[ 2/5][  2574][ 500.45s]    d_loss/g_loss: -0.27/-0.64
[ 2/5][  2674][ 519.71s]    d_loss/g_loss: -0.26/0.17
[ 2/5][  2774][ 538.92s]    d_loss/g_loss: 0.028/ 0.4
[ 3/5][  2811][ 546.22s]    d_loss/g_loss: -0.033/-1.2

```


[3/5][2911][565.48s]	d_loss/g_loss: -0.33/-0.23
[3/5][3011][584.79s]	d_loss/g_loss: -0.026/ 1.0
[3/5][3111][604.07s]	d_loss/g_loss: -0.42/-0.98
[3/5][3211][623.32s]	d_loss/g_loss: 0.14/-0.1
[3/5][3311][642.55s]	d_loss/g_loss: -0.3/ 0.6
[3/5][3411][661.78s]	d_loss/g_loss: -0.093/0.21
[3/5][3511][681.13s]	d_loss/g_loss: -0.1/-1.1
[3/5][3611][700.36s]	d_loss/g_loss: -0.21/-1.2
[3/5][3711][719.63s]	d_loss/g_loss: -0.11/0.16
[4/5][3748][726.92s]	d_loss/g_loss: -0.12/-0.4
[4/5][3848][746.22s]	d_loss/g_loss: -0.045/-0.39
[4/5][3948][765.48s]	d_loss/g_loss: -0.13/-1.0
[4/5][4048][784.76s]	d_loss/g_loss: -0.23/-0.88
[4/5][4148][804.01s]	d_loss/g_loss: -0.2/-1.7
[4/5][4248][823.22s]	d_loss/g_loss: -0.21/-0.6
[4/5][4348][842.48s]	d_loss/g_loss: -0.12/-1.6
[4/5][4448][861.69s]	d_loss/g_loss: -0.1/0.66
[4/5][4548][880.99s]	d_loss/g_loss: -0.0077/-1.8
[4/5][4648][900.24s]	d_loss/g_loss: -0.11/-0.75

✓ Visualization (2 pts)

```

1 # Visualize the loss
2 # include the figure in the latex file (1 pt)
3 plt.title("Generator and Critic Loss During Training")
4 plt.plot(generator_losses,label="G")
5 plt.plot(critic_losses,label="D")
6 plt.xlabel("iterations")
7 plt.ylabel("Loss")
8 plt.legend()
9 plt.show()


```



```

1 # Visualize the generation (you may scroll to see the animation of training)
2 # include the final figure in the latex file (1 pt)
3 import matplotlib.animation as animation
4 from IPython.display import HTML
5 #%%capture
6 fig = plt.figure(figsize=(8,8))
7 plt.axis("off")
8 ims = [[plt.imshow(i.permute(1,2,0), animated=True)] for i in img_list]
9 ani = animation.ArtistAnimation(fig, ims, interval=1000, repeat_delay=1000, blit=True)
10
11 HTML(ani.to_jshtml())

```



☐ Once ☒ Loop ☐ Reflect

